

DATABASE MANAGEMENT SYSTEMS II LAB

CSE4410

SWE 21

CSE
IUT

Contents

Lab 1	Introduction	3
1	Marks Distribution	3
2	Approximate Course Outline	3
3	Task - Group B	4
4	Task - Group A	5
Lab 2	Tablespace	6
1	Default Tablespaces	6
2	Create Tablespace	7
3	Extend Tablespace	7
4	Drop Tablespace	8
5	Read-Only or Read-Write Tablespace	8
6	Online and Offline Tablespace	8
7	Task - Group B	9
8	Task - Group A	9
Lab 3	Java Database Connectivity (JDBC)	10
1	Environment Setup	10
2	Java Database Connectivity with Oracle	10
3	Some interfaces of JDBC API	12
4	Task - Group B	14
5	Task - Group A	15
Lab 4	Functions and Procedures	16
1	Procedure	16
2	Function	17
3	Parameter Notations	18
4	Error Handling	20
5	Task - Group B	22

Lab 1 Introduction

Welcome to CSE 4410.

1 Marks Distribution

Module	Mark (%)
Attendance	10
Lab Evaluation	40
Lab Report	20
Project	30

2 Approximate Course Outline

1. (Intro) + Basics of Relational Database Model
2. Tablespace
3. JDBC Connection + (Project Proposal Submission)
4. PL/SQL
 - a. Function/Procedure
 - b. Cursor
 - c. Trigger
5. Project Progress Presentation
6. NoSQL [MongoDB]
 - a. Theory
 - b. Sessional
7. Graph-based Database [Neo4j]
 - a. Theory
 - b. Sessional
8. Project Presentation

3 Task - Group B

Consider the schema shown in Figure 1.1 for the database of a university:

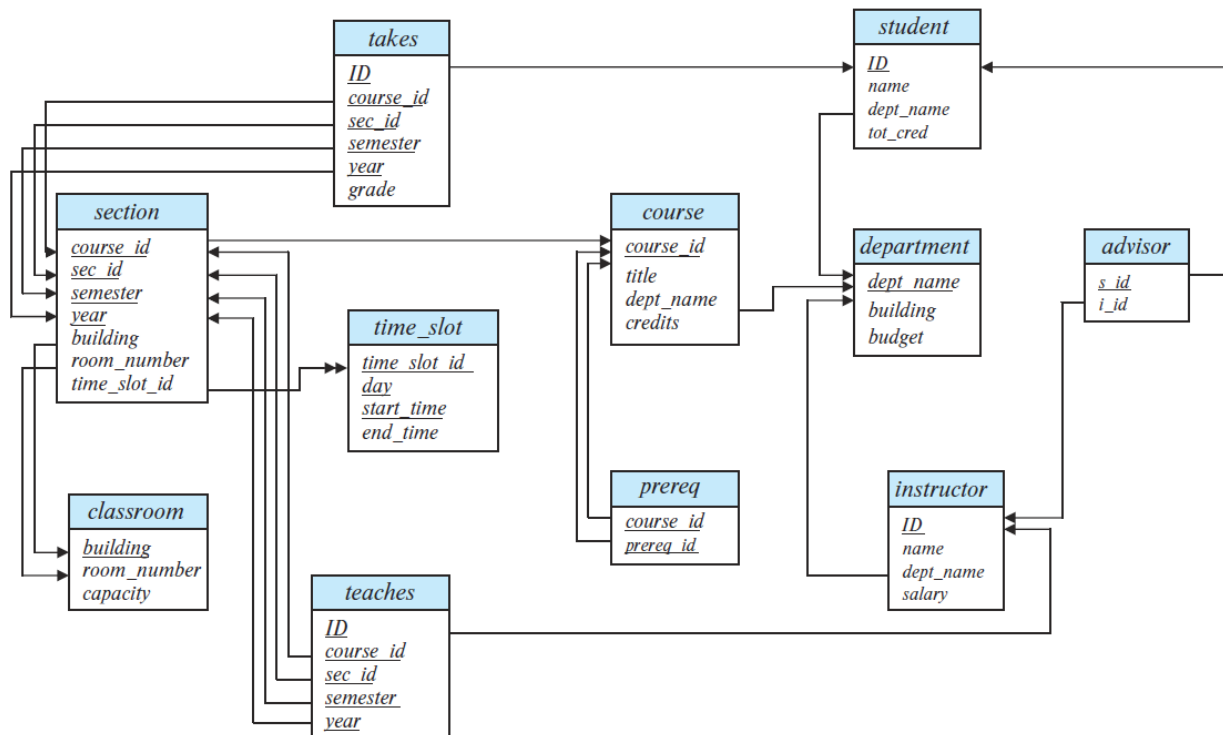


Figure 1.1. Schema diagram for a university database

Write the command @“<file_path>\<file_name>.sql” in your SQL command line to execute the provided .sql files. Now, write SQL statements to answer each of the following queries:

- Find the names of all the instructors from the 'Biology' department.
- Show the Course ID and the Title of all the courses registered for by the student with ID '12345'.
- Find the names and department names of all the students who have taken a course offered by the 'Comp. Sci.' department.
- Find the names of the students who take the 'CS-101' course in 'Spring, 2018'.
- Find the names of students who have taken the highest number of courses with a specific prefix 'CS'.
- Find the names of students who have taken courses taught by at least three different instructors
- Find the course name and section having the minimum number of enrollments. Do not include the sections that do not have any students enrolled.
- Find the name of the instructor, dept_name, and count of students he/she advising. If an instructor is not advising any student, show 0.
- Find the name and department of the students who take more courses than the average number of courses taken by a student.
- Insert each instructor as a student with total credit set to 0 in the same department they are teaching.
- Remove all the newly added students from the previous query.
- Update the 'tot_cred' for each student based on the credits taken.
- Update the salary of each instructor to 10000 times the number of course sections they have taught.
- Grades are mapped to a grade point as follows: A:10, B:8, C:6, D:4, and F:0. Create a table to store these mappings, and write a query to find the Credit Point Information (CPI) of each student, using this table. Make sure students who have not got a non-null grade in any course are displayed with a CPI of null.

4 Task - Group A

Consider the schema shown in Figure 1.2 for the database of a university:

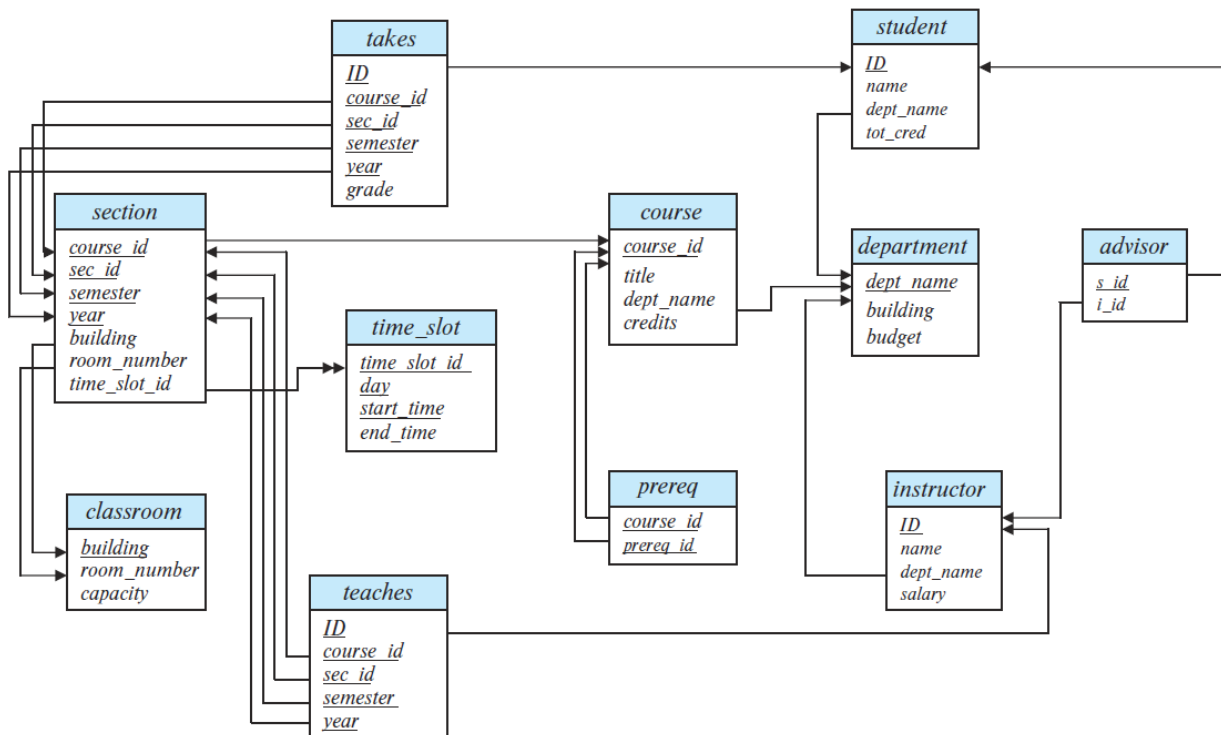


Figure 1.2. Schema diagram for a university database

Write the command @“<file_path>\<file_name>.sql” in your SQL command line to execute the provided .sql files. Now, write SQL statements to answer each of the following queries:

1. Find the names of courses offered by the 'Comp. Sci.' department which has 3 credits.
2. For each student, list their ID, name, and total credits s/he has taken. Do not include the students who did not register for any course.
3. Find the names and the department names of all instructors who have not taught a course.
4. Find all the course titles that do not have any prerequisites.
5. Find the name of the student who takes 2nd, 3rd, and 5th maximum total credits.
6. Find the names of the instructors who are taking courses with no students enrolled. Also, show the name of the courses.
7. Retrieve the course titles and the percentage of students who earned an 'A' grade in each course.
8. Find the number of instructors who have taught the same course in consecutive years.
9. Insert each student as a student with total credit set to 0 in the same department they are teaching.
10. Update the 'tot_cred' for each student based on the credits taken.
11. Update the salary of each instructor to 10000 times the number of course sections they have taught.
12. Find all rooms that have been assigned to more than one section at the same time.
13. Create a view that will show the instructor-wise time slot for 'Fall, 2017' sorted by the instructor_ID, course_ID, section_ID (Instructor_ID, name, his/her course information, section_ID, count of students in that section for the course, and time_slot).

Lab 2 Tablespace

A tablespace is a database storage unit that groups related logical structures together. The database data files are stored in tablespaces. A data file physically stores the data objects of the database such as tables and indices on disk. Using multiple tablespaces allows more flexibility in performing database operations.

- ▶ To enhance performance, segregate user data from data dictionary data. This minimizes conflicts over I/O resources.
- ▶ Prevent cross-application impact by keeping the data of each application separate. This safeguards against disruptions if a tablespace needs to be temporarily disabled.
- ▶ Enhance efficiency by storing data files from different tablespaces on different disk drives. This reduces contention for I/O resources.
- ▶ Increase overall system availability by selectively taking individual tablespaces offline while keeping others operational. This minimizes disruptions.
- ▶ Tailor tablespace allocation to specific database needs, such as high update activity, read-only activity, or temporary segment storage. This optimizes overall database performance.
- ▶ Enhance data security and recovery capabilities by conducting backups at the tablespace level. This facilitates targeted recovery and maintenance efforts.

Some operating systems set a limit on the number of files that can be opened simultaneously (For example, 512 in Windows, 1024 in Ubuntu, 4096 in CentOS, etc.). Efficient tablespace planning is essential to avoid surpassing the operating system limit. It is advisable to create tablespaces based on actual needs, keeping the number of tablespaces to a minimum. When expanding a tablespace, rather than creating numerous small data files, adding one or two substantial data files or opting for data files with autoextension enabled is recommended.

1 Default Tablespaces

Oracle comes with 5 default tablespaces:

- ▶ The primary tablespace in any database is the **SYSTEM** tablespace, which contains information basic to the functioning of the database server, such as the data dictionary and the system rollback segment. The **SYSTEM** tablespace is the first tablespace created at database creation. It is managed as any other tablespace but requires a higher level of privilege and is restricted in some ways. For example, it is not possible to rename or drop the **SYSTEM** tablespace or take it offline.
- ▶ The **SYSAUX** tablespace, which acts as an auxiliary tablespace to **SYSTEM** tablespace, is also created during database creation. It contains the schemas used by various Oracle products and features so that those products do not require their own tablespace. The management of the **SYSAUX** tablespace is similar to that of the **SYSTEM** tablespace.
- ▶ **USERS** is a permanent tablespace containing the application data. Oracle fills this space with the data created and entered by the users.
- ▶ **UNDOTBS1** is an auto-extending tablespace containing the undo data. Oracle provides a fully automated mechanism, referred to as automatic undo management, for managing undo information and space. With automatic undo management, the database manages undo segments in an undo tablespace.
- ▶ **TEMP** is the temporary tablespace that is used for storing intermediate results. Oracle uses it as work areas for tasks such as sort operations for users and sorting during index creation. Oracle does not allow users to create objects in a temporary tablespace. By definition, the temporary tablespace holds data only for the duration of a user's session, and the data can be shared by all users.

2 Create Tablespace

Using **CREATE TABLESPACE** statement, a new tablespace can be created. As we have seen a tablespace consisting of one or more data files, we need to specify the path of the data files as well as their size.

```
1 CREATE TABLESPACE TBS1
2     DATAFILE 'tbs1_data.dbf' SIZE 1M,
3     'tbs2_data.dbf' SIZE 1M;
```

To allow extent management to be **LOCAL**, one can optionally add the statement **EXTENT MANAGEMENT LOCAL AUTOALLOCATE** or **EXTENT MANAGEMENT LOCAL UNIFORM SIZE size**. Locally Managed Tablespaces (LMT) have a bitmap of the blocks or groups of blocks, allowing them to track extent allocation without reference to the data dictionary. If **UNIFORM** is specified, all extents within the tablespace will be the same size, with **1M** (which stands for 1MB) being the default extent size. The **AUTOALLOCATE** clause allows you to size the initial extent leaving Oracle to determine the optimum size for subsequent extents, with 64K being the minimum.

Once the tablespace is created, all the information about it is available in the **DBA_DATA_FILES** view.

```
1 SELECT TABLESPACE_NAME, FILE_NAME, BYTES/1024/1024 MB
2     FROM DBA_DATA_FILES;
```

To assign a user to a specific tablespace, one can explicitly mention it at the time of user creation.

```
1 CREATE USER iutlearner
2     IDENTIFIED BY test123
3     DEFAULT TABLESPACE TBS1;
```

Usually, when we create any new table in Oracle, by default, that is placed in the **USERS** tablespace. However, to create a new table in a user-defined tablespace, one must add the name of the tablespace at the end of the **CREATE TABLE** statement.

```
1 CREATE TABLE T1
2 (
3     ID INT,
4     C1 VARCHAR2(32)
5 ) TABLESPACE TBS1;
```

Then, if we want to check the free space of a certain tablespace, we can fetch that data from the **DBA_FREE_SPACE** view.

```
1 SELECT TABLESPACE_NAME, BYTES/1024/1024 MB
2     FROM DBA_FREE_SPACE
3     WHERE TABLESPACE_NAME='TBS1';
```

3 Extend Tablespace

Commonly, the tablespaces of the database get completely occupied. In that case, no further addition of data is possible. We have already learned about locally managed extent by the time on tablespace creation. But even if we forget to do that or do not want to automate that we can manually handle it using **ALTER TABLESPACE** statement. There are two ways of extension:

► By adding a new data file.

```
1 ALTER TABLESPACE TBS1
2     ADD DATAFILE 'tbs3_data.dbf' SIZE 1M;
```

Here, if we use the **AUTOEXTEND ON** clause at the end of the code, Oracle will automatically extend the size of the data file as needed.

► By resizing the data file

```
1 ALTER DATABASE
2 DATAFILE 'tbs1_data.dbf' RESIZE 15M;
```

4 Drop Tablespace

To remove a tablespace from the database, we use **DROP TABLESPACE** statement.

```
1 DROP TABLESPACE TBS1
2 [INCLUDING CONTENTS [AND | KEEP] DATAFILES]
3 [CASCADE CONSTRAINTS];
```

Here **INCLUDING CONTENTS** is necessary when there is any table created in the tablespace. Any attempt to remove a tablespace that has objects without specifying the clause will result in an error. If we do not use **AND DATAFILES**, it will by default keep the data files of those tables stored without any tablespace. And **CASCADE CONSTRAINTS** is necessary in case of referential integrity.

You can use the **DROP TABLESPACE** command to remove a tablespace regardless of whether it is online or offline. However, it is good practice to take the tablespace offline before removing it to ensure that no sessions are currently accessing any objects in the tablespace.

5 Read-Only or Read-Write Tablespace

The read-only tablespaces allow Oracle to avoid performing edits on large, static parts of a database. It allows you to remove objects such as tables and indexes from a read-only tablespace. However, it does not allow you to create or alter objects in a read-only tablespace.

```
1 ALTER TABLESPACE TBS1 READ ONLY;
2
3 ALTER TABLESPACE TBS1 READ WRITE;
```

By default, any newly created tablespace is in read-write mode.

6 Online and Offline Tablespace

Lastly, a tablespace can be online or offline. If a tablespace is offline, one cannot access data stored in it. On the other hand, if a tablespace is online, its data is available for reading and writing.

```
1 ALTER TABLESPACE TBS1 OFFLINE;
2
3 ALTER TABLESPACE TBS1 ONLINE;
```

Normally, a tablespace is online so that its data is available to users. However, we can take a tablespace offline to make data inaccessible to users when we update and maintain the applications.

7 Task - Group B

1. Create two tablespaces `tbs1` and `tbs2`.
2. Set quota for a single user on both tablespaces.
3. Create two tables `student(name, ID, fk[dept_ID])` and `department(ID, name)` in `tbs1`.
4. Create another table `course(course_code, name, credit, fk[offered_by_dept_ID])` in `tbs2`.
5. Insert a large amount of data in the `student` table and `course` table.
6. List the title and the name of the offering department of each course.
7. Check the free space of the tablespaces.
8. Extend `tbs1` by adding extra data files.
9. Extend `tbs2` by resizing data files.
10. Check the size of the tablespaces.
11. Set `tbs1` to offline and show that the data cannot be accessed.
12. Delete tablespace `tbs1` including the data files.
13. Delete tablespace `tbs2` excluding the data files.

8 Task - Group A

1. Create two tablespaces `tsp1` and `tsp2`.
2. Set quota for a single user on both tablespaces.
3. Create two tables `author(name, ID)` and `publisher(ID, name)` in `tsp1`.
4. Create another table `book(ISBN_code, name, price, fk[publisher_ID])` in `tsp2`.
5. Insert a large amount of data in the `book` table and `publisher` table.
6. List the title and the name of the publisher of each book.
7. Check the free space of the tablespaces.
8. Extend `tsp1` by adding extra data files.
9. Extend `tsp2` by resizing data files.
10. Check the size of the tablespaces.
11. Set `tsp1` to offline and show that the data cannot be accessed.
12. Delete tablespace `tsp1` including the data files.
13. Delete tablespace `tsp2` excluding the data files.

Lab 3 Java Database Connectivity (JDBC)

Java Database Connectivity (JDBC) is a Java API facilitating the interaction with relational databases. It provides a standard interface for connecting to databases, executing SQL queries, manipulating data, and handling the results. Prior to JDBC, ODBC was used but had dependencies on C-language drivers. Later JDBC introduced its own API, employing Java-based JDBC drivers.

The JDBC driver serves as a bridge, allowing the Java application to send queries, receive results, and generally interact with the database management system seamlessly. The driver essentially translates Java calls into a format that the database understands and vice versa, enabling effective and standardized communication between the Java application and the underlying database.

There are four types of drivers. The **JDBC-ODBC bridge driver**, the **Native-API driver** (partially Java driver), the **Network Protocol driver** (fully Java driver), and the **Thin driver** (fully Java driver). Among these, the Thin driver excels in performance, as it directly translates JDBC calls into the specific protocol of the associated database and does not require any supplementary software installation on either the client or server side.

1 Environment Setup

All the files mentioned below have been provided in Google Classroom.

1. Install Java Development Kit (JDK) using `jdk-19_windows-x64_bin.msi`.
2. If you want to use VSCode install the extension pack for Java.
3. Create a new Java Project and add `ojdbc14.jar` or `ojdbc6.jar` file as an external JAR file.
 - a. For VSCode:
from the Explorer Bar, go to **Java Projects** → `<projectname>` → **Referenced Libraries**. Then click on the **plus** sign and select the `ojdbc14.jar` or `ojdbc6.jar` and click the **Select Jar Libraries** button.
 - b. For Eclipse IDE:
from the Menu Bar, go to **Project** → **Properties** →. Then on the opened window, click on **Java Build Path** → **Classpath** → **Add External JARs**. This will open a File Explorer where you can navigate to the path where `ojdbc14.jar` or `ojdbc6.jar` is located to select and add it to your project.
 - c. For IntelliJ IDE:
from the Top Bar, go to **File** → **Project Structure** →. Then on the opened window, click on **Modules**. In the **Export** section, click on the **plus** sign → **JARs or Directories**. There provide the path to `ojdbc14.jar` or `ojdbc6.jar` and hit the **OK** button.

If you are using any other IDE, Google “how to add external JAR files to `<XYZ>` IDE” to get help. Replace `<XYZ>` with the IDE that you are using.

2 Java Database Connectivity with Oracle

There are 5 steps to connect any Java application with the database using JDBC. These steps are as follows:

1. Register the Driver class
2. Create a connection
3. Create statement
4. Execute query
5. Close connection

Consider the Code Snippet 3.1 for connecting with JDBC step by step.

```

1 import java.sql.*;
2
3 class jdbc_practice {
4     public static void main(String args[]) {
5         try {
6             // step1 load the driver class
7             Class.forName("oracle.jdbc.driver.OracleDriver");
8
9             // step2 create the connection object
10            Connection con = DriverManager.getConnection(
11                "jdbc:oracle:thin:@localhost:1521:xe", "DBMS", "dbms");
12
13            // step3 create the statement object
14            Statement stmt = con.createStatement();
15
16            // step4 execute query
17            // drop table
18            System.out.println("Drop table...");
19            String sql = "DROP TABLE REGISTRATION";
20            stmt.executeUpdate(sql);
21
22            // create table
23            String sql1 = "CREATE TABLE REGISTRATION " +
24                "(id INTEGER not NULL, " +
25                " first VARCHAR(255), " +
26                " last VARCHAR(255), " +
27                " age INTEGER, " +
28                " PRIMARY KEY ( id ))";
29
30            stmt.executeUpdate(sql1);
31
32            // insert table
33            System.out.println("Inserting records into the table...");
34            String sql2 = "INSERT INTO Registration VALUES (100, 'Zara', 'Ali',
18)";
35            stmt.executeUpdate(sql2);
36            sql2 = "INSERT INTO Registration VALUES (101, 'Mahnaz', 'Fatma', 25)";
37            stmt.executeUpdate(sql2);
38
39            // select table
40            System.out.println("Selecting records from the table...");
41            String QUERY = "SELECT id, first, last, age FROM Registration";
42            ResultSet rs = stmt.executeQuery(QUERY);
43            while (rs.next()) {
44                // Display values

```

```

45         System.out.print("ID: " + rs.getInt("id"));
46         System.out.print(", Age: " + rs.getInt("age"));
47         System.out.print(", First: " + rs.getString("first"));
48         System.out.println(", Last: " + rs.getString("last"));
49     }
50     rs.close();
51
52     // update table
53     System.out.println("Updating records from the table...");
54     String sql3 = "UPDATE Registration " +
55         "SET age = 30 WHERE id in (100, 101)";
56     stmt.executeUpdate(sql3);
57
58     // delete table
59     System.out.println("deleting records from the table...");
60     String sql4 = "DELETE FROM Registration " +
61         "WHERE id = 101";
62     stmt.executeUpdate(sql4);
63
64     // step5 close the connection object
65     con.close();
66
67     } catch (Exception e) {
68         System.out.println(e);
69     }
70
71 }
72 }

```

Code Snippet 3.1. Sample JDBC Connection code

Consider the following:

- ▶ Driver class: The driver class for the Oracle Database is “`oracle.jdbc.driver.OracleDriver`”.
- ▶ Connection URL: The connection URL for the Oracle10G database is “`jdbc:oracle:thin:@localhost:1521:xe`” where `jdbc` is the API, `oracle` is the database, `thin` is the driver, `localhost` is the server name on which oracle is running (we may also use an IP address here), `1521` is the port number, and `XE` is the Oracle Service Name. You may get all these information from the `tnsnames.ora` file.
- ▶ Username: The default username for the Oracle Database is `system`.
- ▶ Password: It is the password given by the user at the time of installing the Oracle Database.

3 Some interfaces of JDBC API

The `java.sql` package contains classes and interfaces for JDBC API. Some popular interfaces and classes in the JDBC API include:

- ▶ **DriverManager:** `java.sql.DriverManager` is a class that manages a list of database drivers. It is used to establish a connection to the database by selecting an appropriate driver from the list.
- ▶ **Connection:** `java.sql.Connection` represents a connection to a relational database. It provides methods

for creating statements, committing or rolling back transactions, and managing other aspects of the connection.

- **Statement:** `java.sql.Statement` is an interface that represents an SQL statement. There are different types of statements, such as `Statement`, `PreparedStatement`, and `CallableStatement`, each serving a specific purpose. These are used to execute SQL queries and updates.
- **ResultSet:** `java.sql.ResultSet` represents the result set of a database query. It provides methods for iterating over the rows of the result set and retrieving data from each column.
- **PreparedStatement:** `java.sql.PreparedStatement` is a sub-interface of `Statement`. It is used to execute pre-compiled SQL queries with parameters. `PreparedStatement` is more efficient than `Statement` for executing queries repeatedly with different parameter values.
- **ResultSetMetaData:** `java.sql.ResultSetMetaData` is an interface that provides information about the columns of a `ResultSet`, such as the column names, types, and properties.
- **DatabaseMetaData:** `java.sql.DatabaseMetaData` provides methods to obtain metadata about the database, such as information about its tables, columns, and supported SQL features.

4 Task - Group B

Consider the schema shown in Figure 3.1 for the database of a university:

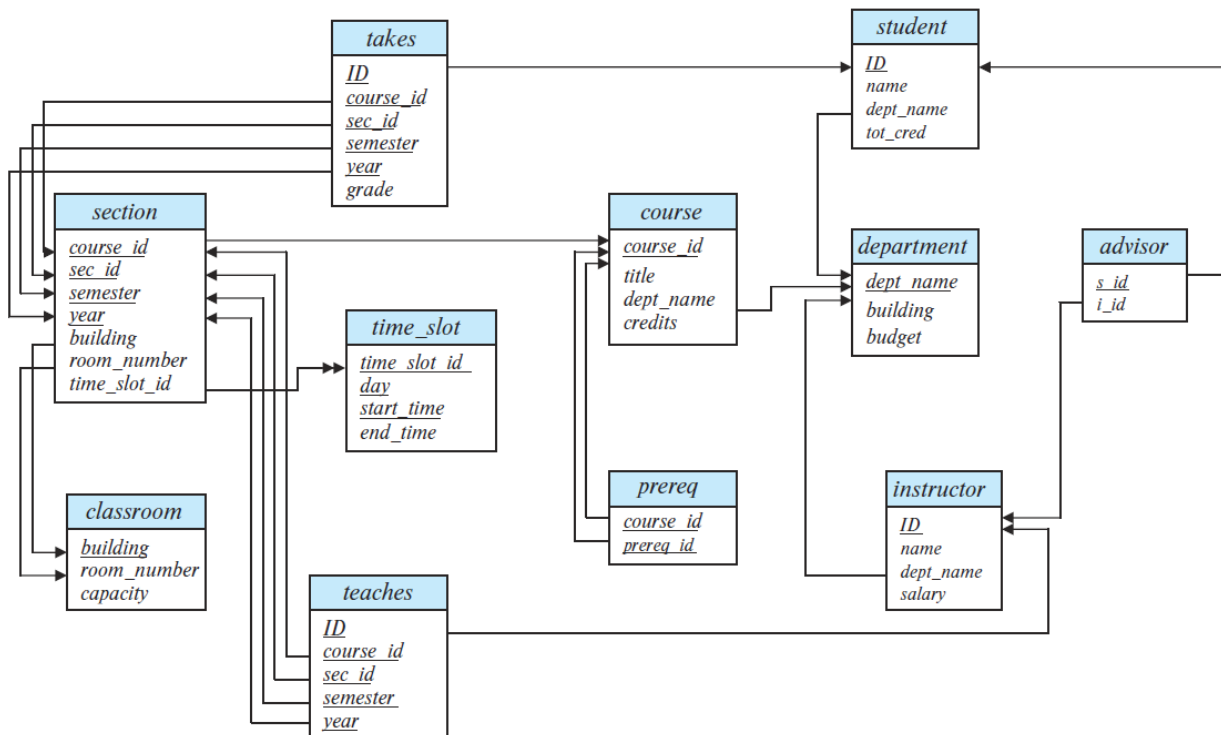


Figure 3.1. Schema diagram for a university database

Write the command @“<file_path>\<file_name>.sql” in your SQL command line to execute the provided .sql files. Now, write Java functions to perform each of the following tasks:

1. Decrease the budget of the departments having a budget of more than 99999 by 10%. Then show the number of departments that did not get affected.
2. Take the day of the week, starting hour, and ending hour as input from the user. Then print the names of the instructors who will be taking classes during that time.
3. Find the top N students based on the number of courses they are enrolled in. You should take N as input and print the ID, name, department name, and the number of courses taken by the student. If N is larger than the total number of students, print the information for all the students.
4. Insert a new student named 'Jane Doe' in the STUDENT table. The student should be enrolled in the department having the lowest number of students. The ID of the student will be $(X + 1)$, where X is the highest ID value among the existing students.
5. Find out the list of students who do not have any advisor assigned to them. Then assign them an advisor from their department. In case there are multiple instructors from a certain department, the advisor should be selected based on the least number of students advised. Finally, print the name of the students, the name of their advisor, and the number of students advised by the said advisor.

5 Task - Group A

Consider the schema shown in Figure 3.2 for the database of a university:

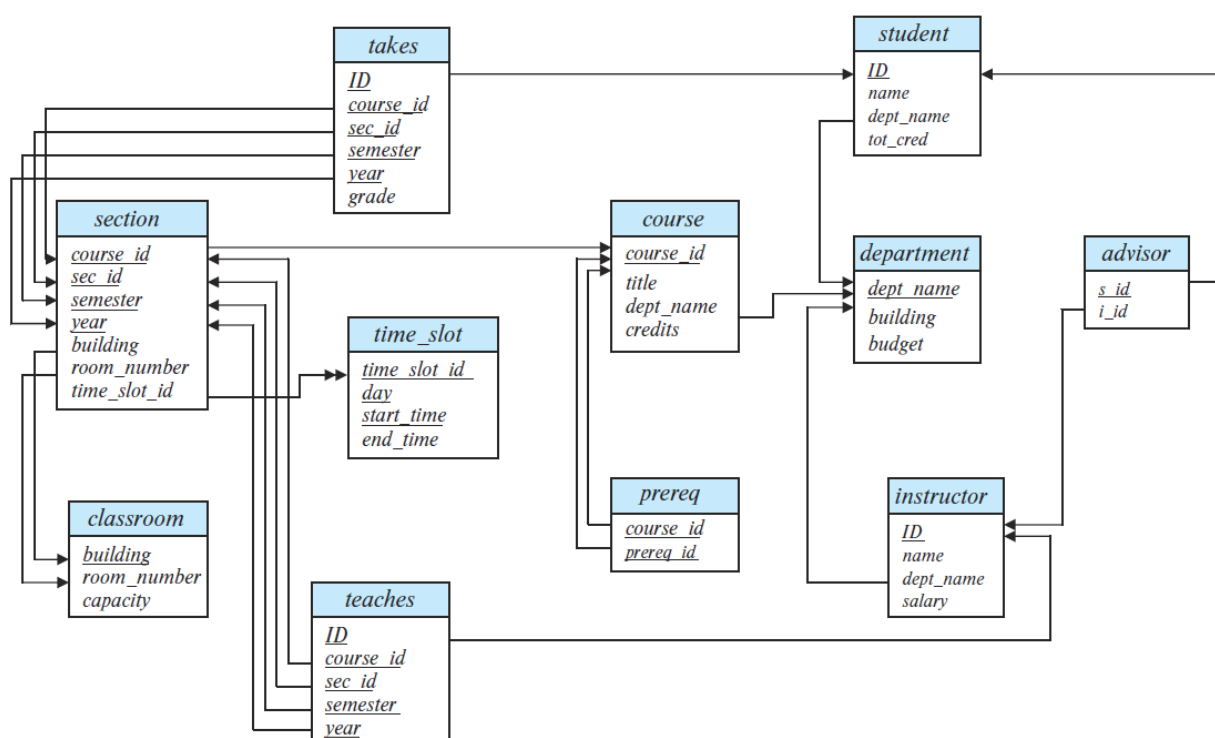


Figure 3.2. Schema diagram for a university database

Write the command @“<file_path>\<file_name>.sql” in your SQL command line to execute the provided .sql files. Now, write Java functions to perform each of the following tasks:

1. Set the salary of each instructor to 9000X where X is the total credits of the courses taught by the instructor. Print the names of the instructors whose salaries remain unchanged.
2. Considering the pre-requisite(s) for each course, print the course title and the names of the students who can enroll in them.
3. Take the name of the student as input from the user. Then print the weekly routine for the student. Each class should be printed in the following format:

```
<Day>
<Start Time> - <End Time>
<Course ID> - <Title>
<Building> - <Room>
```

The days should be sorted based on the regular order of weekdays starting from Monday. If there are multiple classes scheduled on the same day, they should be sorted based on the starting time.

4. Find out the list of instructors who do not have any students assigned to them. Then assign them students from the same department who do not have any advisor. If there are multiple students from the same department that meet the criteria, then select the one with the lowest tot_cred. After that, print the names of the instructors who still do not have any students assigned to them.
5. Insert a new instructor named 'John Doe' in the INSTRUCTOR table. The instructor should be enrolled in the department having the highest number of students. The ID of the instructor will be (X - 1), where X is the lowest ID value among the existing instructors. The salary of the instructor should be the average among all the instructors of the same department. Finally, print the information of the instructor.

Lab 4 Functions and Procedures

Functions and Procedures are two of the key components of PL/SQL stored programming units. They are used to build database-tier libraries to encapsulate application functionality, which is then co-located on the database tier for efficiency. Oracle maintains a unique list of stored object names for tables, views, sequences, stored programs, and types. This list is known as a namespace. Functions and procedures are in this namespace.

Stored functions and procedures provide a way to hide implementation details in a program unit. They also let you wrap the implementation from prying eyes on the server tier. However, the main motivation for Functions and Procedures is modular code. **Modularization is the process by which you break up large blocks of code into smaller pieces** (modules) that can be called by other modules. Modularization of code is analogous to normalization of data, with many of the same benefits and a few additional advantages. With modularization, your code becomes:

- ▶ **Reusable:** By breaking up a large program or entire application into individual components that plug-and-play together, you will usually find that many modules are used by more than one other program in your current application. Designed properly, these utility programs could even be of use in other applications!
- ▶ **Manageable:** Which would you rather debug: a 1,000-line program or five individual 200-line programs that call each other as needed? Our minds work better when we can focus on smaller tasks. You can also test and debug on a per program scale (called unit testing) before individual modules are combined for a more complicated integration test.
- ▶ **Readable:** Modules have names, and names describe behavior. The more you move or hide your code behind a programmatic interface, the easier it is to read and understand what that program is doing. Modularization helps you focus on the big picture rather than on the individual executable statements. You might even end up with that most elusive kind of software: self-documenting code.
- ▶ **Reliable:** The code you produce will have fewer errors. The errors you do find will be easier to fix because they will be isolated within a module. In addition, your code will be easier to maintain because there is less of it and it is more readable.

1 Procedure

A procedure is a module that performs one or more actions. Because a **procedure call is a standalone executable statement in PL/SQL**, a PL/SQL block could consist of nothing more than a single call to a procedure. Procedures are key building blocks of modular code, allowing you to both consolidate and reuse your program logic. The following illustrates a procedure prototype:

```
1 [CREATE [OR REPLACE]]
2 PROCEDURE procedure_name[(parameter[, parameter]...)]
3 [AUTHID {DEFINER | CURRENT_USER}] {IS | AS}
4 [PRAGMA AUTONOMOUS_TRANSACTION;]
5     [local_declarations]
6 BEGIN
7     executable_statements
8 [EXCEPTION
9     exception_handlers]
10 END [procedure_name];
```

Procedures cannot be right operands. Nor can you use them in SQL statements. You move data into and out of PL/SQL stored procedures through their formal parameter list. Here, parameter modes define the behavior of formal parameters. These parameter modes offer you the ability to use pass-by-value or pass-by-reference formal parameters. The three parameter modes: **IN** (default), **OUT**, and **IN OUT** can be used with any subprogram:

- **IN:** The value of the actual parameter is passed into the procedure when the procedure is invoked. Inside the procedure, the formal parameter acts like a PL/SQL constant. **It is considered read-only**, i.e., it cannot be changed. You can assign a default value to a variable with **IN** parameter mode, making it optional.
- **OUT:** Any value the actual parameter has when the procedure is called is ignored. Inside the procedure, the formal parameter acts like an uninitialized PL/SQL variable and thus has a value of **NULL**. It can be read from and written to. Note that, you cannot assign a default value to a variable with **OUT** parameter, as it would make the variable optional.
- **IN OUT:** This mode is a combination of **IN** and **OUT**.

During the declaration of the parameters in Procedures, you must leave out the constraining part of the declaration. The following procedure can be used to determine the salary of an instructor given their ID:

```

1  /* Create procedure */
2  CREATE OR REPLACE
3  PROCEDURE FIND_SAL(I_ID IN NUMBER, SALARY OUT NUMBER)
4  AS
5  BEGIN
6      SELECT MAX(SALARY) INTO SALARY
7          FROM INSTRUCTOR
8          WHERE ID = I_ID;
9  END;
10 /
11
12 /* Call it from an anonymous block */
13 DECLARE
14     AMOUNT NUMBER;
15 BEGIN
16     FIND_SAL(10101, AMOUNT);
17     DBMS_OUTPUT.PUT_LINE(AMOUNT);
18 END;
19 /

```

Output:

```
65000
```

2 Function

A function is a module that returns data through its **RETURN** clause, rather than in an **OUT** or **IN OUT** parameter. Unlike a procedure call, which is a standalone executable statement, a call to a function can exist only as part of an executable statement, such as an element in an expression or the value assigned as the default in a declaration of a variable. Functions are convenient structures because you can call them directly from SQL statements or PL/SQL programs. They can also be used as right operands because they return a value. Since Functions return explicit values, it is recommended to not use **OUT** and **IN OUT** modes with functions. The following illustrates a function prototype:

```

1  [CREATE [OR REPLACE]]
2  FUNCTION function_name[(parameter[, parameter]...)]
3  RETURN return_type
4  [AUTHID {DEFINER | CURRENT_USER}] {IS | AS}
5  [PRAGMA AUTONOMOUS_TRANSACTION;]
6  [local_declarations]

```

```

7 BEGIN
8     executable_statements
9 [EXCEPTION
10     exception_handlers]
11 RETURN statement;
12 END [function_name];

```

Both procedures and functions have a name, can take parameters, return values, and be called by many users. They are stored in the data dictionary. The difference is that a function must return a value, but in a procedure it is optional. Also, you cannot call procedures from an SQL statement.

The following program can be used to calculate the compound interest of a loan:

```

1  /* Declare function */
2  CREATE OR REPLACE
3  FUNCTION COMPOUND_INTEREST(PA NUMBER, AIR NUMBER := 5, TF)
4  RETURN NUMBER
5  IS
6      CI NUMBER;
7  BEGIN
8      CI := PA * ((1 + (AIR / 100)) ** TF) - PA;
9  RETURN CI;
10 END;
11 /
12
13 /* Call it from an anonymous block */
14 BEGIN
15     DBMS_OUTPUT.PUT_LINE(COMPOUND_INTEREST(10000, 5, 5));
16 END;
17 /

```

Output:

```
2762.815625
```

Here, PA denotes the principle amount, AIR denotes the annual interest rate (which has a default value of 5), TF denotes the compound period (years), and CI denotes the compound interest.

3 Parameter Notations

When calling a subroutine, such as a procedure or a function, positional, named, and mixed notations can be used.

3.1 Positional Notation

Positional notation means that you provide a value for each variable in the formal parameter list. The values must be in sequential order and must also match the datatype. You use positional notation to call the functions as follows:

```

1 BEGIN
2     DBMS_OUTPUT.PUT_LINE(COMPOUND_INTEREST(10000, 5, 5));
3 END;
4 /

```

3.2 Named Notation

Named notation means that you pass actual parameters by using their formal parameter name, the association operator (\Rightarrow), and the value. You call a function using named notation by:

```

1 BEGIN
2     DBMS_OUTPUT.PUT_LINE(COMPOUND_INTEREST(TF => 5, PA => 10000, AIR => 5));
3 END;
4 /

```

Named notation lets you only pass values to required parameters, which means you accept the default values for any optional parameters.

```

1 BEGIN
2     DBMS_OUTPUT.PUT_LINE(COMPOUND_INTEREST(TF => 5, PA => 10000));
3 END;
4 /

```

3.3 Mixed Notation

Mixed notation means that you can call subroutines by a combination of positional and named notation. This becomes very handy when parameter lists are defined with all mandatory parameters first and optional parameters next. It lets you name or avoid naming the mandatory parameters, and it lets you skip optional parameters where their default values work. It does not solve exclusionary notation problems. Exclusionary problems occur with positional notation when optional parameters are interspersed with mandatory parameters, and when you call some but not all optional parameters.

```

1 BEGIN
2     DBMS_OUTPUT.PUT_LINE(COMPOUND_INTEREST(10000, TF => 5, AIR => 5));
3 END;
4 /

```

There is a restriction on mixed notation. All positional notation of actual parameters must occur first and in the same order as they are defined by the function signature.

3.4 Exclusionary Notation

If the formal parameters are defined as optional, you can exclude one or more of the actual parameters. For example, consider the following function:

```

1 CREATE OR REPLACE
2 FUNCTION ADD_THREE_NUMBERS
3 (A NUMBER := 0, B NUMBER := 0, C NUMBER := 0)
4 RETURN NUMBER
5 IS
6     SUM NUMBER;
7 BEGIN
8     SUM := A + B + C;
9 RETURN SUM;
10 END;
11 /

```

Here, all 3 parameters are optional. So we can write programs like:

```
1 BEGIN
2     DBMS_OUTPUT.PUT_LINE(ADD_THREE_NUMBERS(3, C => 4));
3 END;
4 /
```

7

When you opt to not provide an actual parameter, it acts as if you are passing a null value. This is known as exclusionary notation. Oracle recommends that you should list optional parameters last in function and procedure signatures. They also recommend that you sequence optional variables so that you never have to skip an optional parameter in the list. These recommendations exist to circumvent errors when making positional notation calls.

You cannot really skip an optional parameter in a positional notation call. This is true because all positional calls are in sequence by datatype, but you can provide a comma-delimited null value when you want to skip an optional parameter in the list. However, Oracle 11g now lets you use mixed notation calls. You can now use positional notation for your list of mandatory parameters, and named notation for optional parameters. This lets you skip optional parameters without naming all parameters explicitly.

4 Error Handling

When creating a procedure/function, you might face compilation errors. For example, if you execute the following code:

```
1 CREATE OR REPLACE
2 PROCEDURE FIND_SAL(I_ID IN NUMBER, SALARY OUT NUMBER)
3 AS
4 BEGIN
5     SELECT MAX(SALARY) INTO SALARY
6     FROM INSTRUCTOR
7     WHERE ID = I_ID /* ERROR: Missing semicolon */
8 END;
9 /
```

Output:

Warning: Procedure created with compilation errors.

Very helpful(!) You can use `SHOW ERROR` command (or `SHO ERR` for short) right after the warning message to find out the errors:

```
1 SHOW ERROR
```

Output:

Errors for PROCEDURE FIND_SAL:

LINE/COL ERROR

```
-----  
4/5      PL/SQL: SQL Statement ignored  
6/29     PL/SQL: ORA-00933: SQL command not properly ended  
7/4      PLS-00103: Encountered the symbol "end-of-file" when expecting  
          one of the following:  
          ( begin case declare end exception exit for goto if loop mod  
          null pragma raise return select update while with  
          <an identifier> <a double-quoted delimited-identifier>  
          <a bind variable> << continue close current delete fetch lock  
          insert open rollback savepoint set sql execute commit forall  
          merge pipe purge
```

This should help you identify that you missed a semicolon (;) in line 7.

5 Task - Group B

Write PL/SQL statements to perform the following tasks:

1. Warm-up

- Print your first name.
- Take your student ID as input and print its length.

2. Consider the schema shown in Figure 4.1 for the database of a university:

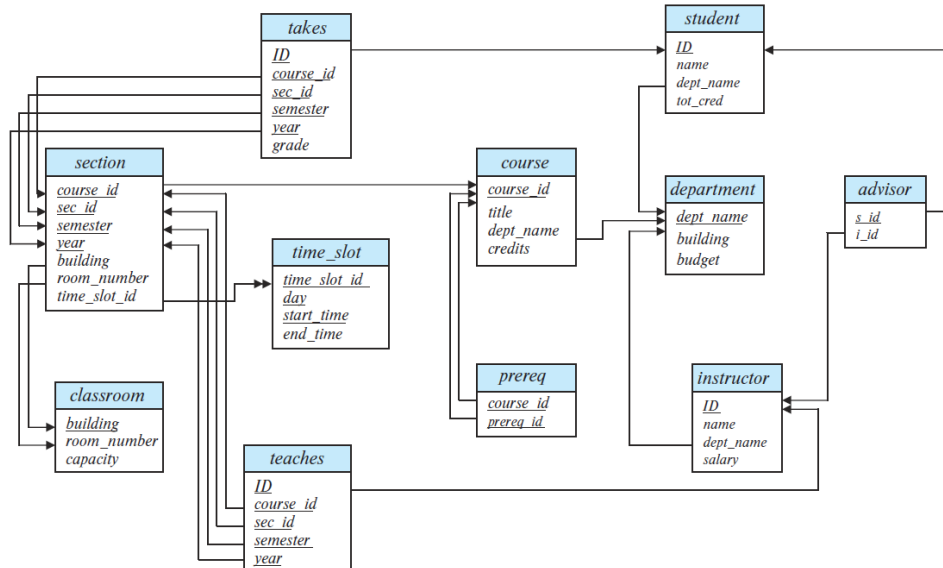


Figure 4.1. Schema diagram for a university database

Write the command @“<file_path>\<file_name>.sql” in your SQL command line to execute the provided .sql files. Now, write Java functions to perform each of the following tasks:

- Set the salary of each instructor to 9000X where X is the total credits of the courses taught by the instructor. Print the names of the instructors whose salaries remain unchanged.
- Considering the pre-requisite(s) for each course, print the course title and the names of the students who can enroll in them.
- Take the name of the student as input from the user. Then print the weekly routine for the student. Each class should be printed in the following format:

```
<Day>
<Start Time> - <End Time>
<Course ID> - <Title>
<Building> - <Room>
```

The days should be sorted based on the regular order of weekdays starting from Monday. If there are multiple classes scheduled on the same day, they should be sorted based on the starting time.

- Find out the list of instructors who do not have any students assigned to them. Then assign them students from the same department who do not have any advisor. If there are multiple students from the same department that meet the criteria, then select the one with the lowest tot_cred. After that, print the names of the instructors who still do not have any students assigned to them.
- Insert a new instructor named 'John Doe' in the INSTRUCTOR table. The instructor should be enrolled in the department having the highest number of students. The ID of the instructor will be (X - 1), where X is the lowest ID value among the existing instructors. The salary of the instructor should be the average among all the instructors of the same department. Finally, print the information of the instructor.